

②

$$a) \min_{\alpha} \sum_{i=1}^N |y_i - \sum_{j=1}^{K-1} \alpha_j b_j(x^{(i)})|^2 + p \int_{x_{\min}}^{x_{\max}} \left(\sum_{j=1}^{K-1} \alpha_j b_j(z) \right)''^2 dz$$

In HW 1, I showed $\sum_{i=1}^N |y_i - \sum_{j=1}^D \alpha_j h_j(x^{(i)})|^2 = \|\gamma - B\alpha\|^2$

where $B \in \mathbb{R}^{N \times D}$. So $\sum_{i=1}^N |y_i - \sum_{j=1}^{K-1} \alpha_j b_j(x^{(i)})|^2 = \|\gamma - B\alpha\|^2$

Where $B \in \mathbb{R}^{N \times (K-1)}$. It remains to show $p \int_{x_{\min}}^{x_{\max}} \left(\sum_{j=1}^{K-1} \alpha_j b_j(z) \right)''^2 dz$

$= p\alpha^T \Omega \alpha$:

$$p \int_{x_{\min}}^{x_{\max}} \left(\sum_{j=1}^{K-1} \alpha_j b_j(z) \right)'' \left(\sum_{j=1}^{K-1} \alpha_j b_j(z) \right)'' dz = p \int_{x_{\min}}^{x_{\max}} \left(\sum_{j=1}^{K-1} \alpha_j b_j''(z) \right) \left(\sum_{i=1}^{K-1} \alpha_i b_i''(z) \right) dz$$

$$= p \int_{x_{\min}}^{x_{\max}} \left(\sum_{j=1}^{K-1} \sum_{i=1}^{K-1} \alpha_j \alpha_i b_j''(z) b_i''(z) \right) dz = p \sum_{j=1}^{K-1} \sum_{i=1}^{K-1} \alpha_j \alpha_i \int_{x_{\min}}^{x_{\max}} b_j''(z) b_i''(z) dz$$

Now for each i, j , $\int_{x_{\min}}^{x_{\max}} b_j''(z) b_i''(z) dz$ is a number, denote it Ω_{ji} (*)

Then we have $p \sum_{j=1}^{K-1} \sum_{i=1}^{K-1} \alpha_j \alpha_i \Omega_{ji}$. By HW 1, we know this

is just $p\alpha \cdot (\Omega\alpha) = p\alpha^T \Omega \alpha$, where $\Omega \in \mathbb{R}^{(K-1) \times (K-1)}$

And so putting everything together we get: $\min_{\alpha} \|\gamma - B\alpha\|^2 + p\alpha^T \Omega \alpha$

where $B \in \mathbb{R}^{N \times (K-1)}$, $\Omega \in \mathbb{R}^{(K-1) \times (K-1)}$ and symmetric and each entry in

Ω is $\Omega_{kl} = \int_{x_{\min}}^{x_{\max}} b_k''(z) b_l''(z) dz$ as shown in (*).

Now let's find the soln to the min problem:

$$\min_{\alpha} [\|\gamma - B\alpha\|^2 + p\alpha^T \Omega \alpha] = \min_{\alpha} [(\gamma - B\alpha)^T (\gamma - B\alpha) + p\alpha^T \Omega \alpha]$$

$$= \min_{\alpha} [\gamma^T \gamma - 2\gamma^T B\alpha + \alpha^T B^T B \alpha + p\alpha^T \Omega \alpha] = \min_{\alpha} f(\alpha)$$

$f(\alpha)$ is quadratic and we know how to take its gradient:

$$\nabla f(\alpha) = -2B^T \gamma + 2B^T B \alpha + 2p\Omega \alpha \quad \text{since both } B^T B \text{ and } \Omega \text{ are symmetric}$$

Solving for the critical point: $\nabla f(\alpha) = 0 \Rightarrow$

$B^T B \alpha + p\Omega \alpha = B^T \gamma \Rightarrow \alpha^* = (B^T B + p\Omega)^{-1} B^T \gamma$, if $(B^T B + p\Omega)$ is invertible ($p > 0$). If $p > 0$, then α^* is a min

HW12

Ines Pancorbo

4/13/2020

2 b)

```
data <- read.csv("BoneMassData.txt", header = TRUE, sep = ' ', stringsAsFactors = TRUE)
```

```
# only working with female
```

```
data_f <- data[which(data$gender == "female"),]
```

```
# 1000 knots equally spaced between
```

```
# the minimum and maximum x values (ages) of data_f
```

```
x_min <- range(data_f$age)[1]
```

```
x_max <- range(data_f$age)[2]
```

```
knots <- seq(x_min, x_max, length.out = 1000)
```

```
x_min
```

```
## [1] 9.4
```

```
x_max
```

```
## [1] 25.55
```

```
# 100000 equally spaced points in mygrid,
```

```
# should be dense enough since range of age is not too big
```

```
mygrid <- seq(x_min, x_max, length.out = 100000)
```

```
# distance between points
```

```
h <- mygrid[2] - mygrid[1]
```

```
library("splines")
```

```
B <- splineDesign(knots = knots,  
                  x = data_f$age[order(data_f$age)],  
                  outer.ok = T)
```

```
Bpp <- splineDesign(knots = knots,  
                    x = mygrid,  
                    derivs = 2,  
                    outer.ok = T)
```

```
# checking dim of Bpp, should be 100000 x (1000-4)
```

```
dim(Bpp)
```

```
## [1] 100000    996
```

```
# checking dim of B,
```

```
# should be (length(data_f$age)=259) x (1000-4)
```

```
dim(B)
```

```
## [1] 259 996
```

The Ω matrix via numerical integration is:

```
# Approximating each entry in Omega matrix using Riemann integration
omega <- (t(Bpp) %*% Bpp)*h
```

```
# double-checking dimensions of omega, should be (1000-4) x (1000-4)
dim(omega)
```

```
## [1] 996 996
```

```
# printing the first 20 entries of first row of omega
omega[1,1:20]
```

```
## [1] 6.311774e+05 -3.550285e+05 -5.518051e+00 3.944960e+04 0.000000e+00
## [6] 0.000000e+00 0.000000e+00 0.000000e+00 0.000000e+00 0.000000e+00
## [11] 0.000000e+00 0.000000e+00 0.000000e+00 0.000000e+00 0.000000e+00
## [16] 0.000000e+00 0.000000e+00 0.000000e+00 0.000000e+00 0.000000e+00
```

2 c)

For any given B , $B^T B$ is symmetric and positive semi-definite. Positive semi-definite because for any given x , $x^T B^T B x = (Bx)^T (Bx) = \|Bx\|^2 \geq 0$. Because $B^T B$ is positive semi-definite, we are not guaranteed $B^T B$ is invertible. For my given B , lets see if $B^T B$ is invertible or not.

```
# Lets first show B^TB is not invertible
det(t(B) %*% B)
```

```
## [1] 0
```

The determinant of $B^T B$ is 0 so $B^T B$ is not invertible.

Now, Ω is symmetric since by definition each entry of Ω is $\Omega_{kl} = \int_{x_{min}}^{x_{max}} b_k''(z)b_l''(z)dz$ and multiplication is commutative (i.e., $b_k''(z)b_l''(z) = b_l''(z)b_k''(z) \Rightarrow \Omega_{kl} = \Omega_{lk}$). Further, it is positive semi-definite by construction: $x^T \Omega x = \int_{x_{min}}^{x_{max}} ((\sum_{j=1}^{996} x_j b_j(z))'')^2 dz \geq 0$. Again, we are not guaranteed invertibility but lets see if Ω is invertible.

```
# Using the determinant to check invertability is not
# the best approach
# You can easily get an overflow during the numerical computation of the determinant
# For example:
det(omega)
```

```
## [1] Inf
```

```
# The fact that det = infty might indicate omega is close to singular
# but R is still able to give you an inverse:
solve(omega)[1,1:5]
```

```
## [1] 6.765061e-06 1.169034e-05 1.708152e-05 2.232080e-05 2.757378e-05
```

```
# Given numerical precision, a better approach is finding kappa(matrix),
# if this value exceeds 10^16, like we saw in class, then the matrix
# is computationally singular
kappa(omega)
```

```
## [1] 49872460901
```

```
log10(kappa(omega))
```

```
## [1] 10.69786
```

So $\text{kappa}(\omega)$ is approx $10^{11} < 10^{16}$, and therefore ω is computationally invertible (\Rightarrow invertible). Ω is consequently positive definite.

Now it is easy to show that for my given B (which is positive semi-definite), for my given Ω (which is positive definite), and for any given $\rho > 0$, $B^T B + \rho \Omega$ will be positive definite and thus, invertible:

If $\rho > 0$ then $\rho \Omega$ is still positive definite: By the spectral decomposition theorem, $\Omega = Q D Q^T \Rightarrow \rho \Omega = \rho(Q D Q^T) \Rightarrow \rho \Omega = Q \rho D Q^T$. In other words $\rho \Omega$ is a symmetric matrix with the same eigenvectors as Ω and eigenvalues equal to the eigenvalues of Ω times ρ (consequently, all positive eigenvalues).

Now, since by definition of $\rho \Omega$ being positive definite we have $x^T \rho \Omega x > 0$ and by definition of $B^T B$ being positive semi-definite we have $x^T B^T B x \geq 0$, we have as a result that for any given x , $x^T (B^T B + \rho \Omega) x = x^T B^T B x + x^T \rho \Omega x > 0$.

Consequently, we have that for my given B , for my given Ω , and for any given $\rho > 0$, $B^T B + \rho \Omega$ will be positive definite and thus, invertible.

2 d)

```
# function
smooth_spline_regression <- function(rho){

  # calculating the coefficients of the spline given work in 2 a)
  alpha <- solve(t(B)%*%B+rho*omega,
                t(B)%*(data_f$spnbmd[order(data_f$age)]))


  ##### spline #####
  spline <- B%*%alpha
  #####

  # plot
  plot(data_f$age, data_f$spnbmd, col = "grey")
  points(data_f$age[order(data_f$age)],
         spline,
         col = "blue",
         lwd = 2,
         type = "l")
  title(paste("Smooth Spline Regression with rho = ", rho))
}
```


Now, using the function defined above, calculate alpha using $\rho = 0.01, 1, 100$ and plot resulting spline.

```
for (rho in c(0.01, 1, 100)){
  smooth_spline_regression(rho)
}
```


Smooth Spline Regression with $\rho = 0.01$



Smooth Spline Regression with $\rho = 1$



Smooth Spline Regression with $\rho = 100$



Comment on the fit in each case:

It looks like for $\rho = 0.01$ the spline regression is overfitting our data (the wiggly line, and how this wiggly line passes through some of the data points) and so a penalty term of $\rho = 0.01$ might be too small. For $\rho = 100$, it looks like the spline regression is underfitting our data (it looks almost linear) so a penalty term of $\rho = 100$ might be too big. Lastly, for $\rho = 1$, the spline regression seems like a decent fit (it looks like it is capturing the “shape” of the data, and the spline isn’t too “wiggly” or too linear).

3 a)

```
# reading in as matrix A
A <- as.matrix(read.csv("A.txt", header = FALSE, sep=" "))
first_row_A <- A[1,]

# reading in as matrix no_noise_A
no_noise_A <- as.matrix(read.csv("no_noise_A.txt", header = FALSE, sep=" "))
first_row_no_noise_A <- no_noise_A[1,]


# getting signal
sig <- seq(1,2,length.out = 10)
sig[5] <- -1

# getting q
q <- as.matrix(read.csv("q.txt", header = FALSE, sep=" "))
```

Now, lets plot the values in the first row of A, the first row of no_noise_A, and the underlying signal.


```
library(ggplot2)
ggplot() +
  geom_line(aes(x = seq(1,ncol(A)), y = first_row_A, color = "First row of A")) +
```

```
geom_line(aes(x = seq(1,ncol(A)), y = first_row_no_noise_A, color = "First row of no_noise_A")) +
geom_line(aes(x = seq(1,ncol(A)), y = sig, color = "Signal")) +
xlab("") + ylab("")
```



Now, let's plot the values in the first row of A , the 10 values corresponding to the means of the 10 columns of A , and the signal to see if it is possible to figure out the signal from the first row of A or the $\text{colMeans}(A)$.

```
library(ggplot2)
ggplot() +
  geom_line(aes(x = seq(1,ncol(A)), y = first_row_A, color = "First row of A")) +
  geom_line(aes(x = seq(1,ncol(A)), y = sig, color = "Signal")) +
  geom_line(aes(x = seq(1,ncol(A)), y = colMeans(A), color = "Mean of Columns of A")) +
  xlab("") + ylab("")
```



At least from the plot, it doesn't look like there is a relationship/connection between the column means of A and the signal, or the first row of A and the signal. So I would say, no, I cannot tell what the signal is by looking at the means of the 10 columns of A , or by looking at the first row of A .

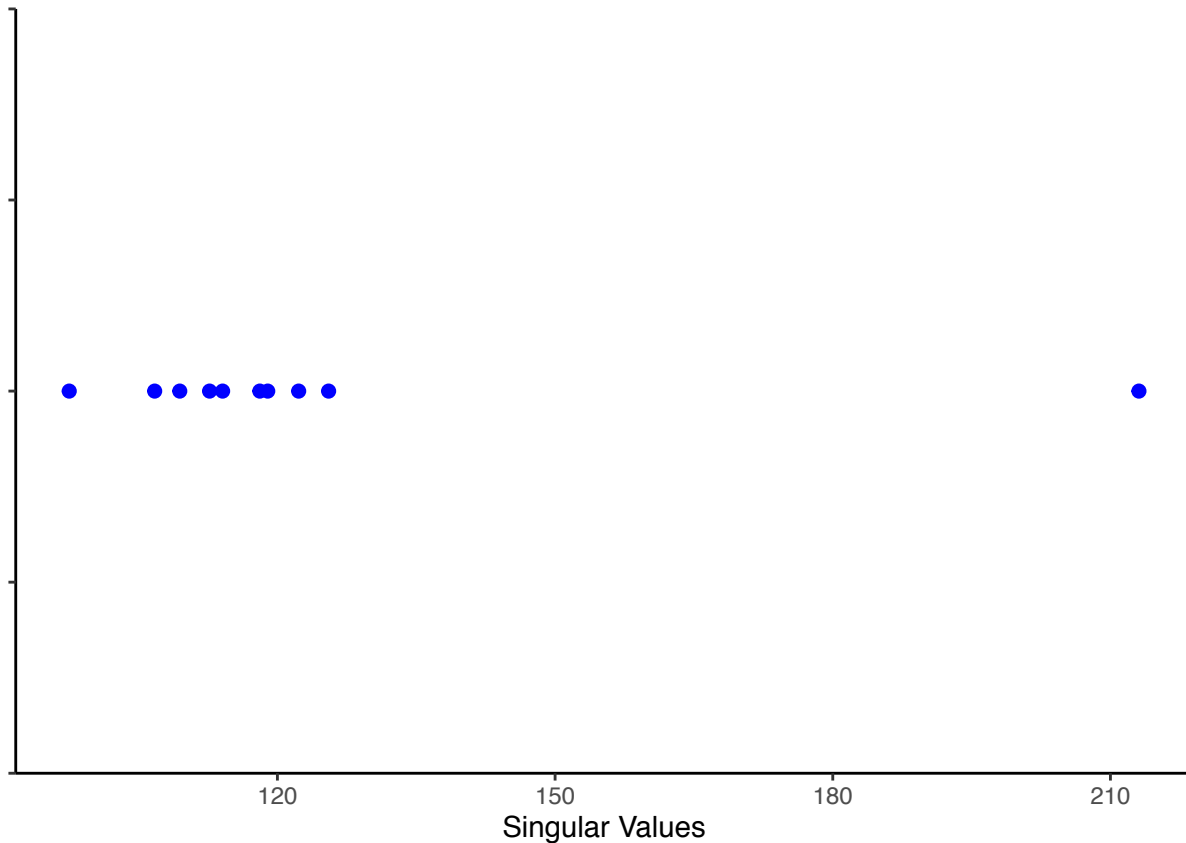
3 b)

```
# find svd of A
svd <- svd(A)
# get U and V
U <- svd$u
V <- svd$v

# get singular values of A
singular_values <- svd$d
# should get 10
length(singular_values)

## [1] 10

# plot singular values of A
ggplot() +
  geom_point(aes(x=singular_values, y = rep(0,length(singular_values))), size=2, color='blue') +
  theme_classic() +
  theme(axis.text.y=element_blank()) +
  xlab("Singular Values") + ylab("")
```

Comment on the singular values:

Some things to keep in mind first:

- (1) The rank of a matrix can represent the amount of unique information stored by a matrix (higher rank => more unique information) and it is the number of non-zero singular values
- (2) One can think of the SVD as a method that decomposes A into a linear combination of rank-1 matrices weighted by the corresponding singular value


Having mentioned the above: the SVD of matrix A will express matrix A as a sum of 10 rank-1 matrices weighted by the corresponding singular value. From the plot of singular values, you can see that there is one that dominates. Therefore, the plot is essentially telling us that we can approximate A well with the rank-1 matrix corresponding to that dominant singular value (i.e., approximate A with $s_1(u_1 v_1^T)$, where s_1 is the dominant singular value, entry (1,1) of the diagonal matrix, and u_1, v_1 are the first column vectors of the matrices U and V). In other words, we can think of matrix A (rank 10) as a rank 1 matrix. This makes sense given that matrix A (rank 10) is just a “noisy” version of matrix `no_noise_A` (rank 1).

Lets now find this rank-1 approximation of A ($s_1(u_1 v_1^T)$) and denote it A_1 .


```
# get matrices U and V
U <- svd$u
V <- svd$v

# find A1
A1 <- singular_values[1]* (U[,1] %*% t(V[,1]))


# Now use image() to visualize the three
# matrices and confirm that A1 removes the noise from A
image(A)
```



```
image(no_noise_A)
```



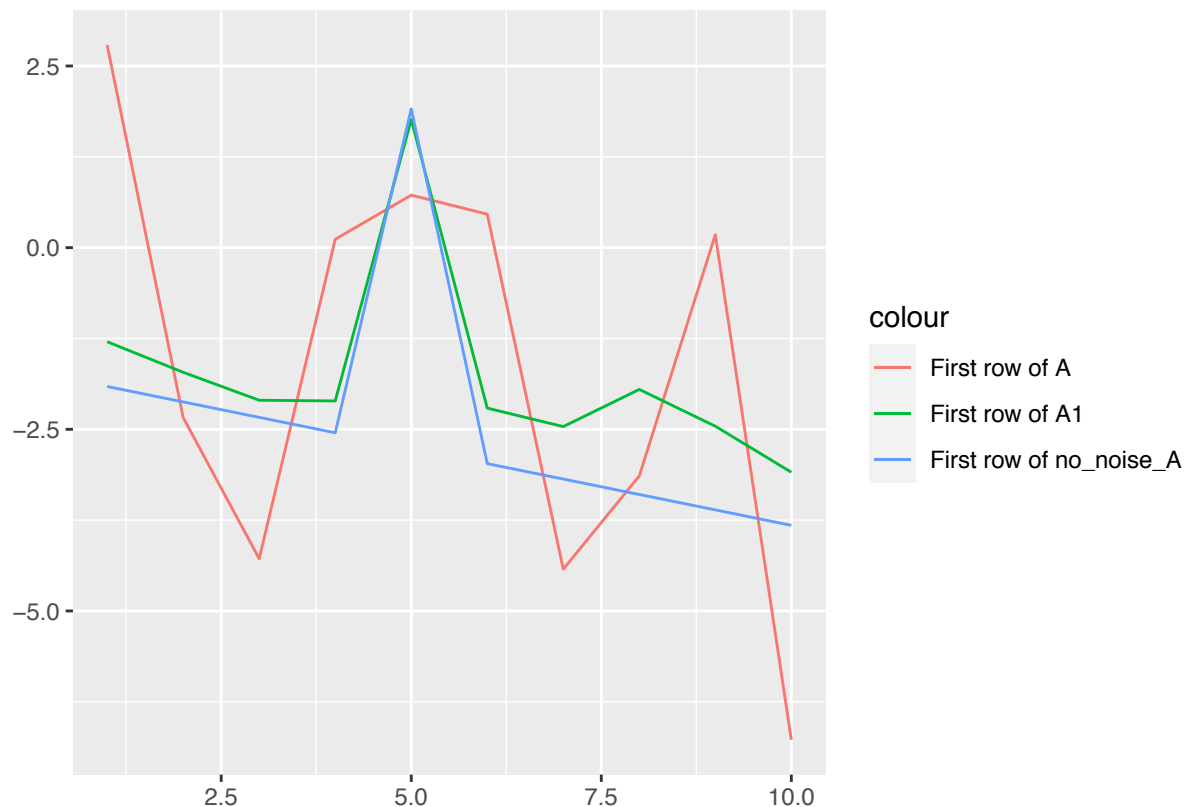
```
image(A1)
```



From what was explained above, it would be reasonable to expect that this rank-1 approximation A_1 of A be similar to `no_noise_A`. That is, we should expect A_1 to throw out much of the “noise” and little of the “underlying signal and vector q .” From the images, since A_1 is very similar to `no_noise_A` (a little bit blurrier but still displays similar patterns) it is reasonable to conclude that A_1 removes the noise from A .

Now let's compare the first rows of A , `no_noise_A`, and A_1 (the 1-rank approx of A corresponding to the dominant singular value) by plotting them.

```
library(ggplot2)
ggplot() +
  geom_line(aes(x = seq(1,ncol(A)), y = first_row_A, color = "First row of A")) +
  geom_line(aes(x = seq(1,ncol(A)), y = A1[1,], color = "First row of A1")) +
  geom_line(aes(x = seq(1,ncol(A)), y = no_noise_A[1,], color = "First row of no_noise_A")) +
  xlab("") + ylab("")
```



From what we saw with `image()`, it makes sense that the first rows of A_1 and `no_noise_A` are similar. The plot above confirms this. The first row of A is not similar to the first row of A_1 or first row of `no_noise_A` due to the presence of noise.

Answering last question: **Given a row of A in the form $q \cdot sig + noise$, what role do v_1 , s_1 and u_1 have in capturing q and sig ?**

Let $U\Sigma V^T$ be the svd of the matrix A . Two things first:

- (1) By construction $A = q(sig)^T + noise$, where q , sig are column vectors in R^{500} and R^{10} respectively, and noise is a 500×10 matrix. So it follows that the i th row of A will be $(q_i)sig + noise[i,]$.
- (2) We know that U is a 500×500 matrix, and V is a 10×10 matrix. Since A has rank 10 we know that the first 10 columns of U are an orthonormal basis for the column space of A , and the 10 columns of V are an orthonormal basis for the row space of A .


From the points discussed above, any row of A will be in the span of V . Lets take the first row of A as an example:

```
# coefficients for linear combination
solve(V,A[1,])
```

```
## [1]  6.8511699  1.1537398 -0.5037349  0.3690536 -4.1620982  0.8170252
## [7]  2.6494192 -4.5955250  1.6988247 -3.1671041
```

So each row of A is connected to each of the 10 columns vectors v_i in V in that sense (each row of A is a linear combination of the v_i 's). However, from the work done above, we can denoise A (i.e., $A \approx q(sig)^T$) by approximating A with the 1-rank matrix $s_1(u_1 v_1^T)$ (i.e., $A \approx s_1(u_1 v_1^T) \approx q(sig)^T$). So the row space of $q(sig)^T$ can be approximated with the span of the vector v_1 . And so it is clear, that v_1 should be able to "capture" the vector sig . If we plot v_1 and sig we should see similarity.

```
ggplot() +
  geom_line(aes(x = seq(1,ncol(A)), y = V[,1], color = "v_1")) +
  geom_line(aes(x = seq(1,ncol(A)), y = sig, color = "signal")) +
  xlab("") + ylab("")
```




Note: the vectors in U and V are defined up to a sign. That is, the svd is not exactly “unique”: One can flip any right singular vector, as long as the corresponding left singular vector is flipped as well. And so, taking this into account, you can see how v_1 (or $-v_1$) captures the signal (both the signal and v_1 (or $-v_1$) lines are shaped similarly, even if scale is different).

Now, that we have discussed the role v_1 plays in any given row of A (it captures the *sig* vector), let's turn to q .

Since by construction, $A = q(\text{sig})^T + \text{noise}$, each column of A is the vector q times a corresponding entry in *sig* plus some *noise*. So let's think of q in terms of the column space of A .

Now, since A is of rank 10, the first 10 column vectors in U are an orthonormal basis for the column space of A . However, from work done above, since approximating A with $s_1(u_1 v_1^T)$ will give you an approximate denoised version of A (i.e., $s_1(u_1 v_1^T) \approx q(\text{sig})^T$), the first vector in U , u_1 , will be an approximate basis for the column space of $q(\text{sig})^T$. So the column space of $q(\text{sig})^T$ can be approximated with the span of the vector u_1 . So it becomes clear that u_1 should be able to “capture” the vector q , and in terms of mappings we can think of $(s_1)u_1$ mapping a scalar to q . Therefore, we should see a relationship between $(s_1)u_1$ and q if we plot them together.

```
ggplot() +
  geom_line(aes(x = seq(1,nrow(A)), y = U[,1]*singular_values[1], color = "s_1(u_1)")) +
  geom_line(aes(x = seq(1,nrow(A)), y = q, color = "vector q")) +
  xlab("") + ylab("")
```



As can be seen from the plot $(s_1)u_1$ captures the vector q (even though scale is different, the line corresponding to q and the line corresponding to $(s_1)u_1$ display a similar shape). We can zoom in to see how they intersect in the range they share:

```
ggplot() +
  geom_line(aes(x = seq(1,nrow(A)), y = U[,1]*singular_values[1], color = "s_1(u_1)")) +
  geom_line(aes(x = seq(1,nrow(A)), y = q, color = "vector q")) +
  xlab("") + ylab("") +
  ylim(-5,5)
```

