

Ines Pancorbo

Kernel Machines

1. Using the dot product for similarity

- $u^T v = ||u|| ||v|| \cos \theta \implies \frac{u^T v}{||u|| ||v||} = \cos \theta \in [-1, 1]$
- We can think u is/is not similar to v in terms of direction, using the above.
- Remember the dot product $u^T v$ is just the projection of vector u onto vector v .

Note: The difference between scalar projection and vector projection.

2. What is a Kernel?

Kernel is a way of computing the dot product of two vectors x and y in some feature space (possibly very high dimensional). Suppose we have a mapping $\phi : R^n \rightarrow R^m$ where $m \gg n$ that brings our vectors in R^n to some feature space R^m . Then the dot product in this space is $\phi(x)^T \phi(y)$. A kernel is a function K that corresponds to this dot product, i.e., $K(x, y) = \phi(x)^T \phi(y)$.

Why would this be useful?

Kernels give a way to compute dot products in some feature space (must be Hilbert space as we need the dot product) without even knowing what the space is and what ϕ is.

For example: Consider a simple polynomial kernel $K(x, y) = (1 + x^T y)^2, x, y \in R^2$. This doesn't seem to correspond to any mapping function ϕ , it is just a function that returns a scalar. Assuming $x = (x_1, x_2)$ and $y = (y_1, y_2)$ let us expand this expression.

$K(x, y) = (1 + x^T y)^2 = 1 + x_1^2 y_1^2 + x_2^2 y_2^2 + 2x_1 y_1 + 2x_2 y_2 + 2x_1 x_2 y_1 y_2$. Note this is nothing but a dot product between two vectors $\phi(x) = (1, x_1^2, x_2^2, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1 x_2)$ and $\phi(y) = (1, y_1^2, y_2^2, \sqrt{2}y_1, \sqrt{2}y_2, \sqrt{2}y_1 y_2)$. So the Kernel $K(x, y) = (1 + x^T y)^2 = \phi(x)^T \phi(y)$ computes a dot product in 6-dimensional space without explicitly visiting this space or passing $x \in R^2$ to $\phi(x) : R^2 \rightarrow R^6$.

Linear Regression using the Kernel Trick

How do we fit the data?

Example:

Data: $(x, y), x \in R^n, y \in R$.

Consider $\phi(x) : R^n \rightarrow R^m$ i.e., $(x_1, x_2, \dots, x_n) \rightarrow (\phi(x_1), \phi(x_2), \dots, \phi(x_n))$.

We want $y \sim f(x) = w^T \phi(x)$, which is m dimensional.

First think of an analogy:

Suppose you wanted a linear model, so $y \sim f(x) = \alpha^T x$ for each x . So the space of functions you can fit from is $F = \{f(x) = \alpha^T x\}$, which is a 3 dimensional space.

We would want to $\min_{\alpha \in R^n} ||y - \beta \alpha||^2$, where β is the model matrix.

$$\beta = \begin{pmatrix} x_1^{(1)} & x_2^{(1)} \\ x_1^{(2)} & x_2^{(2)} \\ x_1^{(3)} & x_2^{(3)} \\ \vdots & \vdots \\ x_1^{(N)} & x_2^{(N)} \end{pmatrix}$$

We know how to solve this minimization problem as it is just a quadratic: $\alpha = (\beta^T \beta)^{-1} \beta^T y$.

F is a small dimensional space (especially when compared to samples N , which we can assume is more than 2) so there would be no need to penalize but if desired our minimization problem and solution would be:

$\min_{\alpha \in R^n} (\|y - \beta \alpha\|^2 + \lambda \|\alpha\|^2)$ (just a ridge regression), with solution $\alpha = (\beta^T \beta + \lambda I)^{-1} \beta^T y$.

Now, going back to a kernel machine fit. We are dealing with a much bigger space in terms of dimensions, $F = \{f(x) = w^T \phi(x)\}$.

We would initially want $\min_{w \in R^n} \|y - \beta w\|^2$, where β is a $N \times m$ matrix

$$\beta = \begin{pmatrix} \phi(x^{(1)})^T \\ \phi(x^{(2)})^T \\ \phi(x^{(3)})^T \\ \vdots \\ \phi(x^{(N)})^T \end{pmatrix}$$

Take the case when $m \gg N$, then the minimization problem we are trying to solve doesn't have a unique solution since we have a fat matrix. And therefore, we have to penalize: $\min_{w \in R^m} (\|y - \beta w\|^2 + \lambda \|w\|^2)$ (again, ridge regression). And the solution is $w = (\beta^T \beta + \lambda I)^{-1} \beta^T y$.

**** Note**, using ridge regression because of quadratic optimization. We can use lasso regression and would be dealing with convex optimization.

However, we have a **problem** from a computational point of view:

β is a $N \times m$ matrix $\implies \beta^T \beta$ is a $m \times m$ matrix. And so computing the inverse, which requires gaussian elimination would require at least $\frac{n^3}{3}$ multiplications. If m is "big," this may be computationally infeasible. So as F grows in dimension (m grows) we require more computation to fit the data.

This is when the kernel trick (i.e., a realization, method) comes into play.

Kernel Trick

Trick #1:

Consider $\beta^T = (\phi(x^{(1)}) \phi(x^{(2)}) \dots \phi(x^{(N)}))$.

We write $w = \beta^T a + z$. We are taking w and expressing it as a linear combination of the feature vectors, $\phi(x)$, and some z . I.e., we are projecting w onto the span of the $\phi(x^{(i)})$, and adding z , which is orthogonal to each of the $\phi(x^{(i)})$.

It turns out that $w \in \text{span}(\phi(x^{(1)}), \phi(x^{(2)}), \dots, \phi(x^{(N)}))$ so that $z = 0$. Let us show why this is true.

We want to $\min_{w \in R^m} (\|y - \beta w\|^2 + \lambda \|w\|^2)$. Now substitute for w , $\beta^T a + z$. Note $a \in R^N, z \in R^m$.

So we have

$$\min_{w \in R^m} \|y - \beta w\|^2 + \lambda \|w\|^2$$

$$\begin{aligned}
&= \min_{a,z} \|y - \beta(\beta^T a + z)\|^2 + \lambda \|\beta^T a + z\|^2 \\
&= \min_{a,z} \|y - (\beta\beta^T a + \beta z)\|^2 + \lambda(\beta^T a + z)^T(\beta^T a + z) \\
&= \min_{a,z} \|y - (\beta\beta^T a + \beta z)\|^2 + \lambda(\beta^T a + z)^T(\beta^T a + z)
\end{aligned}$$

But we have that z is orthogonal to each of the $\phi(x^{(i)})$ so $\beta z = 0$. I.e., I don't need to add a z because in terms of the minimization problem the z has no effect in getting a better fit for y :

$$\begin{aligned}
&\min_{a,z} \|y - (\beta\beta^T a + \beta z)\|^2 + \lambda(\beta^T a + z)^T(\beta^T a + z) \\
&= \min_{a,z} \|y - (\beta\beta^T a + \beta z)\|^2 + \lambda(a^T \beta^T \beta a + 2a^T \beta z + z^T z) \\
&= \min_{a,z} \|y - \beta\beta^T a\|^2 + \lambda(a^T \beta^T \beta a + z^T z) \\
&= \min_{a,z} \|y - \beta\beta^T a\|^2 + \lambda a^T \beta^T \beta a + \lambda z^T z \\
&= \min_a \|y - \beta\beta^T a\|^2 + \lambda a^T \beta^T \beta a
\end{aligned}$$

Because we want to minimize, and thus we should choose $z = 0$.

Summary:

We can make the following substitution $w = \beta^T a$ and the optimization problem becomes $\min_{a \in R^N} \|y - \beta\beta^T a\|^2 + \lambda a^T \beta^T \beta a$. So we go from an optimization problem in R^m to an optimization in R^N even though we are trying to fit $y \sim w^T \phi(x)$ and $w \in R^m$. This means that our above problem is no longer a problem.

Now consider the matrix $\beta\beta^T$:

$$\beta\beta^T = \begin{pmatrix} \phi(x^{(1)})^T \\ \phi(x^{(2)})^T \\ \phi(x^{(3)})^T \\ \vdots \\ \phi(x^{(N)})^T \end{pmatrix} \begin{pmatrix} \phi(x^{(1)}) & \phi(x^{(2)}) & \dots & \phi(x^{(N)}) \end{pmatrix}$$

Trick #2:

The above matrix is the kernel matrix, i.e., $K = \beta\beta^T$. Note that the entry $K_{ij} = \phi(x^{(i)})^T \phi(x^{(j)})$. Now we come to another problem, storing β, β^T will be an issue as well because of the dimension of the $\phi(x^{(i)}) \in R^m$.

Now, the second part of the kernel trick is that we do not have to store the matrices β, β^T . Note the kernel matrix is $N \times N$. To compute a all we need is the dot products of the $\phi(x^{(i)})$, i.e., the Kernel matrix where $K_{ij} = \phi(x^{(i)})^T \phi(x^{(j)})$.

Summary:

data: $(x_{(i)}, y_i), x_{(i)} \in R^n, i = 1, 2, \dots, N$.

We use a feature map, $x \in R^n \rightarrow \phi(x) \in R^m$.

We want $y \sim f(x) = w^T \phi(x)$ where $f \in F$ and $\dim(F)$ "big."

And to find w we want to $\min_{w \in R^m} (\|y - \beta w\|^2 + \lambda \|w\|^2)$. Then I let $w = \beta^T a$ and so what I want is

$\min_{a \in R^N} \|y - \beta\beta^T a\|^2 + \lambda a^T \beta^T \beta a$. So we have a more reasonable optimization. There is a further simplification: we let $K = \beta\beta^T$ and $K_{ij} = \phi(x^{(i)})^T \phi(x^{(j)})$. Thus, $\min_{a \in R^N} \|y - Ka\|^2 + \lambda a^T Ka$. Now, this is a quadratic and

thus the solution is $a = (K^T K + \lambda K^{-1}) K^T y$. This can be further simplified: Note K is symmetric since the dot product is commutative:

$$\begin{aligned} a &= (K^T K + \lambda K^{-1}) K^T y \\ a &= (K(K + \lambda I))^{-1} K y \\ a &= (K + \lambda I)^{-1} K^{-1} K y \\ a &= (K + \lambda I)^{-1} y \end{aligned}$$

Quick checkpoint: Given new x what is predicted $y = w^T \phi(x)$?

Answer:

$$y = w^T \phi(x) = (\beta^T a) \phi(x) = a^T \beta \phi(x) = a^T \begin{bmatrix} \phi(x^{(1)})^T \\ \phi(x^{(2)})^T \\ \phi(x^{(3)})^T \\ \vdots \\ \phi(x^{(N)})^T \end{bmatrix} \phi(x) = a^T \begin{bmatrix} \phi(x^{(1)})^T \phi(x) \\ \phi(x^{(2)})^T \phi(x) \\ \phi(x^{(3)})^T \phi(x) \\ \vdots \\ \phi(x^{(N)})^T \phi(x) \end{bmatrix}$$

Now let us go over two examples:

Example 1:

data: $(x^{(i)}, y_i), x^{(i)} \in R^n, i = 1, 2, \dots, N$.

We use a feature map, $x \in R^n \rightarrow \phi(x) = (1, x_1, x_2, \dots, x_n, x_1^2, x_1 x_2, \dots, x_1 x_n, x_2^2, x_2 x_3, \dots, x_n^2) \in R^{2n+1+n(n-1)/2}$.

$\phi(x)$ can be very high dimensional, but computing $\phi(x^{(i)}) \cdot \phi(x^{(j)})$ is easy.

We can notice that $\phi(x^{(i)}) \cdot \phi(x^{(j)}) = (1 + x^{(i)} \cdot x^{(j)})^2$. Computationally we prefer to do a dot product in R^n than in $R^{2n+1+n(n-1)/2}$. So now we can write dot products for $\phi(x) \in R^{2n+1+n(n-1)/2}$ based on dot products of corresponding $x \in R^n$.

Now, we want $y \sim f(x) = w^T \phi(x)$. We compute K , with $K_{ij} = \phi(x^{(i)}) \cdot \phi(x^{(j)}) = (1 + x^{(i)} \cdot x^{(j)})^2$

$$K = \begin{pmatrix} (1 + x^{(1)} \cdot x^{(1)})^2 & (1 + x^{(1)} \cdot x^{(2)})^2 & \dots & (1 + x^{(1)} \cdot x^{(N)})^2 \\ \vdots & \vdots & \vdots & \vdots \\ (1 + x^{(N)} \cdot x^{(1)})^2 & (1 + x^{(N)} \cdot x^{(2)})^2 & \dots & (1 + x^{(N)} \cdot x^{(N)})^2 \end{pmatrix}$$

Now K will give us $a \in R^N$ used to find $w \in R^{2n+1+n(n-1)/2}$, which is $a = \beta^T w$. I.e., then $a = (K + \lambda I)^{-1} y$. And so $w = \beta^T$. But we **never write w explicitly**. I.e., given new x , $y = w^T \phi(x) = a^T (\beta \phi(x))$ and $\beta \phi(x)$ is also handled through $\phi(x^{(i)}) \cdot \phi(x) = (1 + x^{(i)} \cdot x)^2$.

Comments:

- (1) Applying kernels to logistic regression (explored in HW 15)
- (2) Support Vector Machines (explored in course on Machine Learning)
- (3) Radial Basis Kernel (explored in course on Machine Learning)